

Optimal Portfolio Allocation

WiDS 5.0 Endterm Report

Name: Ashish Shukla

Roll no : 25M1542

Email id : 25m1542@iitb.ac.in

Course: Winter In Data Science (WiDS 5.0)

Institution: IIT Bombay

Date: 31/01/26

Tools used : Python

NumPy

Pandas

Matplotlib

CVXPY

Google Colab

Abstract

This project explores portfolio allocation using modern portfolio theory and computational optimization. The objective was to understand how risk and return interact in financial decision-making and how investor preferences influence optimal portfolio construction. Through a series of assignments and a final optimization project, I learned how to model financial systems mathematically and implement those models using Python. The focus of the work was not data collection but conceptual understanding, experimentation, and optimization. This report summarizes the concepts learned, implementation methods, and insights gained from the WiDS 5.0 program.

Overall, the project combined mathematical reasoning, programming, and visualization to create a practical understanding of portfolio management. It reinforced the idea that finance is deeply connected to data science and optimization, and that structured models can replace intuition-driven decision-making with measurable trade-offs.

1. Introduction

Before starting this project, my understanding of investing was largely informal. I associated financial markets with uncertainty, news cycles, and personal judgment. The WiDS Optimal Portfolio Allocation project introduced a very different perspective. It showed that portfolio management is not guesswork but a structured mathematical problem that can be studied, modelled, and optimized.

At the center of portfolio theory is a simple but powerful question:
how should an investor allocate capital across assets to balance risk and return?

Different investors answer this question differently. Some prioritize safety and predictable growth, while others are willing to accept large fluctuations in exchange for higher expected gains. This diversity of preference means there is no single “best” portfolio. Instead, optimal allocation depends on investor behavior, constraints, and tolerance for uncertainty. Understanding how to represent these preferences mathematically became one of the main goals of the project.

The project progressed step by step from basic financial foundations to advanced optimization. Initially, the focus was on understanding asset behavior — how returns are measured, how volatility represents risk, and how relationships between assets affect portfolio stability. These ideas provided the groundwork for modern portfolio theory. Once the theoretical base was clear, the next stage involved implementing those ideas computationally using Python.

This report documents that learning journey, from foundational concepts to practical implementation and interpretation of results.

3. Concept understanding

The first phase of the project focused on building conceptual clarity. Before implementing any optimization models, it was necessary to understand the financial ideas that make portfolio allocation meaningful. This stage was less about coding and more about developing intuition for how assets behave and how portfolios reduce risk.

One of the most important realizations during this phase was that risk is not a vague idea — it is measurable. Volatility, defined as the standard deviation of returns, provides a numerical way to express uncertainty. Seeing risk as a number rather than a feeling changed how I interpreted financial behavior. Assets that appear attractive in terms of return may carry hidden instability, and the model forces that instability into the decision process.

Another key concept was diversification. through the mathematical framework, I understood why diversification works. The covariance and correlation structure between assets determines how risk spreads across a portfolio. If assets do not move together, combining them creates stability even if each asset individually is risky.

This phase also introduced the idea that portfolio performance is not just about picking strong assets. It is about constructing relationships between assets. The portfolio behaves as a system, and the system's behavior cannot be predicted by looking at components in isolation.

In this phase I learnt about the concepts of expected return, volatility, covariance, and diversification provided the foundation needed for optimization, Which set the foundation for mathematical modelling .

4. Assignment Implementation

Once the foundational concepts were clear, I moved into hands-on assignments where theory became code. This phase was where the project stopped being abstract and started feeling real. Each assignment required implementing financial formulas and verifying them through Python, which forced me to understand every step rather than memorizing equations.

The early assignments focused on calculating returns and risk metrics. I converted stock prices into log returns using NumPy and learned why log returns are preferred in financial modeling. Seeing the transformation directly in code helped me understand that returns are not just numbers — they are the building blocks of statistical modeling.

Another key assignment involved generating covariance and correlation matrices. Before this project, I had never seen how asset relationships could be expressed as matrices. Once implemented, I realized that portfolio behavior depends heavily on how assets move relative to each other. Two risky assets can reduce overall portfolio risk if their movements offset each other. That insight only became clear after seeing the matrix structure in code.

The assignments also trained me in debugging and validation. Financial calculations are sensitive. A small mistake can produce unrealistic results. I learned to verify outputs, interpret numbers logically, and question unexpected behavior. This process strengthened both my programming discipline and my analytical thinking.

By the end of the assignment phase, I was equipped with hands on experience on portfolio optimization.

5. Portfolio Optimization Project

The optimization project marked the turning point of the entire experience. Instead of observing portfolio behavior, I began actively constructing optimal portfolios using mathematical optimization.

I implemented Modern Portfolio Theory using convex optimization. The objective was to balance expected return and risk under realistic constraints. Writing the optimization problem in CVXPY felt surprisingly natural because the syntax mirrors the math. Decision variables represented portfolio weights, the objective function expressed risk-return tradeoffs, and constraints enforced full investment without short selling.

The Efficient Frontier was the most visually powerful result of the project. Generating thousands of portfolios and plotting risk against return produced a clear frontier curve. Seeing this curve emerge from computation was one of the most satisfying moments of the project. It proved that optimization is not guesswork — it is geometry driven by statistics.

I also implemented the Minimum Variance Portfolio. What surprised me was that the safest portfolio was not simply cash or a single low-risk asset. Instead, it was a combination of assets whose correlations balanced each other. Even volatile assets contributed to stability when paired correctly. This showed that diversification is mathematical, not intuitive.

The introduction of the risk aversion parameter λ added a behavioral dimension to the model. By changing λ , I could move along the Efficient Frontier. High λ produced conservative portfolios. Low λ chased return. This experiment demonstrated that portfolio optimization is not about finding a universal best portfolio — it is about matching portfolios to investor psychology.

The optimization project unified everything I learned. Returns, covariance, and diversification were no longer separate concepts. They became inputs to a decision engine that produced rational allocation strategies.

6. Risk Preference Experiment

The final stage of the project extended optimization by explicitly modeling investor behavior. Instead of searching for a single optimal portfolio, I explored how allocation changes when investor risk tolerance varies. This introduced a behavioral dimension to the model and made the results more realistic.

Using a tunable risk-aversion parameter λ , I solved the optimization problem repeatedly. Each value of λ represented a different investor personality. Large values corresponded to conservative investors who prioritize stability, while small values represented aggressive investors willing to accept volatility in exchange for higher expected return.

Plotting portfolio weights against λ revealed a smooth transition in allocation. Assets did not suddenly appear or disappear; their weights adjusted gradually. This continuity showed that portfolio choice is not binary. It exists on a spectrum of tradeoffs.

Another visualization mapped these portfolios onto the Efficient Frontier. This demonstrated how investor preference determines where someone sits on the curve. The frontier does not prescribe one correct answer. It provides a menu of rational options.

This experiment reinforced the idea that finance is not just about numbers. It is about translating human behavior into mathematical structure.

7. Observations

Several insights emerged from the full sequence of assignments and experiments.

First, diversification is far more powerful than intuition suggests. Combining assets with imperfect correlation reduces risk in a measurable way. The model proved that stability can be engineered mathematically.

Second, risk is not the enemy of investing — unmanaged risk is. Optimization does not eliminate risk; it allocates it intelligently. The Efficient Frontier represents the best achievable tradeoffs, not risk-free choices.

Third, portfolio construction is a systems problem. Individual asset quality matters less than asset interaction. Strong assets can produce weak portfolios if they move together. Average assets can produce strong portfolios if they balance each other.

Finally, optimization replaces emotional decision-making with structured reasoning. Instead of reacting to market sentiment, the model produces decisions based on statistical evidence.

These insights changed how I interpret financial systems. Investing is not random. It is governed by relationships that can be modeled and optimized.

8. Learning

Beyond financial theory, this project strengthened my technical skills. Implementing optimization required disciplined programming and careful data handling. I became more comfortable with matrix operations, numerical stability, and debugging complex workflows.

The project also improved my ability to interpret results. Financial outputs are not just numbers; they are decisions encoded in mathematics. Learning to read those outputs critically was as important as computing them.

Perhaps the most valuable outcome was confidence. I moved from seeing optimization as advanced theory to seeing it as a practical tool. The gap between textbook equations and working systems disappeared.

This experience demonstrated that data science and finance are deeply connected. Portfolio theory is a natural application of statistics, linear algebra, and computation.

9. Limitations

While the project provided strong conceptual understanding, it operates under simplified assumptions.

The model assumes ideal conditions. Real investing involves transaction costs, liquidity limits, taxes, and human behavior, which can change portfolio decisions.

The analysis also relies on historical patterns. Markets evolve, and past behavior does not guarantee future performance. The covariance structure used in optimization is an estimate, not a certainty.

Finally, the asset universe was intentionally small to maintain clarity. Real portfolios involve hundreds or thousands of instruments.

These limitations do not invalidate the model, but they highlight the difference between theoretical optimization and real-world complexity.

10. Conclusion

This project fundamentally changed how I understand portfolio allocation. It showed that investing can be approached systematically using mathematics, statistics, and computation.

The transition from intuition-driven thinking to optimization-driven reasoning was the most important shift. Portfolios are not guesses — they are solutions to structured problems.

By combining theoretical foundations with practical implementation, the WiDS project created a complete learning experience. It demonstrated that finance is not separate from data science. It is one of its most powerful applications.

The skills developed through this project extend beyond investing. They represent a framework for making rational decisions under uncertainty — a framework that applies to many fields.